

Digital Career Institute

Python Course - Advanced SQL



Advanced Data Types - Enumerated Types

Types of data that comprise
an ordered set of values.

- Days of the week
Monday < Tuesday < Wednesday < Thursday < ...
- Months
January < February < March < April < May < ...
- Hierarchical position
Intern < Employed < Coordinator < ...
- Stepped processes
To define < To start < Doing < Reviewing < ...

Enumerated

Basic definition of an enumerated type.

```
ENUM(  
    "Selling point",  
    "Local office",  
    "Headquarters"  
);
```

In PostgreSQL a new type has to be defined first:

```
CREATE TYPE location_type AS ENUM(  
    "Selling point",  
    "Local office",  
    "Headquarters"  
);  
ALTER TABLE Location ADD COLUMN type location_type;
```

In other RDBMS the **ENUM** declaration may be used directly in the **ADD COLUMN** clause.

Validation

```
UPDATE "Location" SET type = 'Something else';
```

```
ERROR:  invalid input value for enum location_type:  
"Something else"
```

```
UPDATE "Location" SET type = 'SELLING POINT';
```

```
ERROR:  invalid input value for enum location_type:  
"Something else"
```

Enumerated types are case sensitive.

Sorting

```
UPDATE Location SET type = 'Local office';
UPDATE Location SET type = 'Headquarters'
WHERE name = 'Headquarters';
UPDATE Location SET type = 'Selling point'
WHERE name = 'Location 3';

SELECT * FROM Location
ORDER BY type ASC;
```

The enum fields are sorted according to the order each item was given on the field definition.

Result

name	city_id	type
Location 3	2	Selling point
Location 2	1	Local office
Location 4	3	Local office
Location 5	4	Local office
Location 6	21	Local office
Headquarters	2	Headquarters

Relational

```
SELECT * FROM Location  
WHERE type > 'Selling point';
```

It uses the order of the value in the field definition to evaluate the relational expression.

Result

name	city_id	type
Headquarters	2	Headquarters
Location 2	1	Local office
Location 4	3	Local office
Location 5	4	Local office
Location 6	21	Local office

Advanced Data Types

- UUID

Integer identifiers are good enough for most cases, but sometimes we want our IDs to be unique across different datasets or applications.

Universally Unique Identifiers

a0eebc99-9c0b-4ef8-bb6d-6bb9bd380a11

```
ALTER TABLE Location ADD COLUMN my_uuid_field uuid;
```

Validation

+

Relational

PostgreSQL does not provide functions to generate UUIDs.

The module uuid-oss can be installed for that purpose.

Advanced Data Types - JSON Types

JavaScript Object Notation

```
{ "key1": "Value 1", "key2": 136 }
```

Validation

```
ALTER TABLE Location ADD COLUMN info json;
```

```
UPDATE Location SET info = 23  
WHERE type = 'Headquarters';
```

```
ERROR:  invalid input syntax for type json
```

```
UPDATE Location SET info = '{"forgot": "a_curly"}'  
WHERE type = 'Headquarters';
```

```
ERROR:  invalid input syntax for type json
```

```
UPDATE Location SET info = '{"all": "ok", "none": 2}'  
WHERE type = 'Headquarters';
```

```
Query returned successfully
```

Basic JSON specific operators

`{field} -> {key}`

Returns a JSON data type.

`{field} ->> {key}`

Returns a text data type.

The JSON Type

DLI

Querying

```
SELECT
    info->'all', pg_typeof(info->'all'),
    info->>'all', pg_typeof(info->>'all')
FROM Location
WHERE info->>'all' = 'ok'
```

Result

?column?	pg_typeof	?column?	pg_typeof
"ok"	json	ok	text

The JSON Type

DLI

Querying: nested JSON paths

{field} #> {key}

{field} #>> {key}

```
UPDATE Location SET info = '{"all": {"ok": true}}'
WHERE type = 'Headquarters';

SELECT
    info#>'{all,ok}', pg_typeof(info#>'{all,ok}'),
    info#>>'{all,ok}', pg_typeof(info#>>'{all,ok}')
FROM Location WHERE info#>>'{all,ok}' = 'true';
```

Result

?column?	pg_typeof	?column?	pg_typeof
true	json	true	text

The type **jsonb** is very similar to the type **json** but it is stored in **binary form**.

Advantages over the **json** type:

- more efficient,
- significantly faster to process,
- supports indexing.

The JSONB Type

DLI

Definition

```
ALTER TABLE Location ADD COLUMN infob jsonb;
```

```
UPDATE Location SET infob = '{"all": "ok", "none": 2}'  
WHERE type = 'Headquarters';
```

```
Query returned successfully
```

Additional operators: inclusion

`{json} @> {json}`

Returns **True** if the left JSON includes the right JSON.

`{json} <@ {json}`

Returns **True** if the right JSON includes the left JSON.

The JSONB Type

Additional operators: inclusion

```
SELECT name, infob, infob@>'{"all": "ok"}' FROM Location;
```

name	infob	?column?
Location 2		
Location 4		
Location 5		
Location 6		
Location 3		
Headquarters	{ "all": "ok", "none": 2 }	t

The JSONB Type

DLI

Additional operators: key exists

```
SELECT name, infob, infob?'all', infob?'something'  
FROM Location;
```

{jsonb} ? {key}

Returns **True** if the left
JSON has a key with
the name **key**.

name	infob	?column?	?column?
Location 2			
Location 4			
Location 5			
Location 6			
Location 3			
Headquarters	{ "all": "ok", "none": 2 }	t	f

Advanced Data Types

-

Array Types

A field can be defined as an array of any other data type.

[145, 543, 234]

Arrays

Appending `[]` to a data type will define an array of elements of that type.

```
ALTER TABLE Location  
ADD COLUMN quarterly_earnings integer[];
```

```
ALTER TABLE City  
ADD COLUMN alternate_name varchar[];
```

```
ALTER TABLE Country  
ADD COLUMN boundaries jsonb[];
```


Arrays

```
UPDATE Location
SET quarterly_earnings = ARRAY[0, 0, 0, 0];

UPDATE Location
SET quarterly_earnings = ARRAY[10, 14, 12, 13]
WHERE name = 'Headquarters';

UPDATE Location
SET quarterly_earnings = '{5, 4, 8, 10}'
WHERE name = 'Location 2';

UPDATE Location
SET quarterly_earnings = ARRAY[2, 3, 4, 1]
WHERE name = 'Location 3';

UPDATE Location
SET quarterly_earnings[4] = 3
WHERE name = 'Location 3';
```

name character vary	quarterly_earnings integer[]
Location 4	{0,0,0,0}
Location 5	{0,0,0,0}
Location 6	{0,0,0,0}
Headquarters	{10,14,12,13}
Location 2	{5,4,8,10}
Location 3	{2,3,4,3}

Array indexes start at 1.

Arrays

Accessing elements in the array

```
SELECT
    name,
    type,
    quarterly_earnings[1] AS Q1,
    quarterly_earnings[2] AS Q2
FROM Location
WHERE quarterly_earnings[2] < quarterly_earnings[1];
```

name	type	Q1	Q2
Location 2	Local office	5	4

Multidimensional Arrays

Append as many `[]` as dimensions in the array.

```
ALTER TABLE Location  
ADD COLUMN opening_times integer[][];
```

```
[[8, 12], [13, 17]]
```

Multidimensional Arrays

```
UPDATE Location
SET opening_times = ARRAY[[8, 12], [13, 17]];

UPDATE Location
SET opening_times = '{{11, 20}}'
WHERE type = 'Selling point';

UPDATE Location
SET opening_times[2][2] = 19
WHERE name = 'Headquarters';
```

name character vary	opening_times integer[]
Location 4	{{8,12},{13,17}}
Location 5	{{8,12},{13,17}}
Location 6	{{8,12},{13,17}}
Location 2	{{8,12},{13,17}}
Location 3	{{11,20}}
Headquarters	{{8,12},{13,19}}

Multidimensional Arrays

```
UPDATE Location
SET opening_times = '{{11, 20}, {21}}'
WHERE type = 'Selling point';

ERROR:  malformed array literal: "{{11, 20}, {21}}"
DETAIL:  Multidimensional arrays must have sub-arrays
with matching dimensions.
```

All arrays of the main array
must have the same
length.

Array Operators

DLI

INCLUDES

{array} @> {array}
{array} <@ {array}

CONCATENATE

{array} || {array}

OVERLAPS

{array} && {array}

Array Functions

DLI

APPEND ()

```
array = array_append(  
    array,  
    element  
)
```

REMOVE ()

```
array = array_remove(  
    array,  
    element  
)
```

LENGTH ()

```
array = array_length(  
    array,  
    integer  
)
```

The **integer** indicates the depth of the array hierarchy, starting at 1 as the first dimension of the array.

Advanced Data Types - Time Types

SQL has specific types to manage time data.

`0001-01-01 00:00:00`

Temporal Data

DLI

There are 4 basic types to work with temporal data.

DATE

TIME

INTERVAL

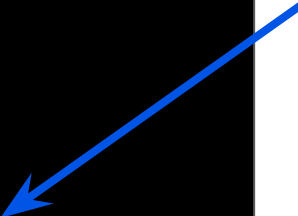
TIMESTAMP

Defining Time Fields

```
ALTER TABLE Location
ADD COLUMN opened_on date;
ALTER TABLE Location
ALTER COLUMN opened_on TYPE time;
ALTER TABLE Location
ALTER COLUMN opened_on TYPE timestamp;

ALTER TABLE Location
ALTER COLUMN opened_on TYPE time with time zone;
ALTER TABLE Location
ALTER COLUMN opened_on TYPE timestamp with time zone;
```

Time and timestamp can also be made aware of the time zone.



Using Time Fields

```
ALTER TABLE Location
ADD COLUMN opened_on date;
UPDATE Location SET opened_on = '1999-01-23';

ALTER TABLE Location
ALTER COLUMN opened_on TYPE time;
UPDATE Location SET opened_on = '14:21:02';

ALTER TABLE Location
ALTER COLUMN opened_on TYPE timestamp;
UPDATE Location SET opened_on = '2004-10-19 10:23:54';
```

Using Time Fields

```
ALTER TABLE Location
ALTER COLUMN opened_on TYPE time with time zone;
UPDATE Location SET opened_on = '14:21:02 PST';

ALTER TABLE Location
ALTER COLUMN opened_on TYPE timestamp with time zone;
UPDATE Location
SET opened_on = '2004-10-19 10:23:54+02';


ALTER TABLE Location
ALTER COLUMN opened_on TYPE time[];
UPDATE Location
SET opened_on = [time '14:21:02', time '15:34:21'];
```

Using Time Fields

```
SELECT * FROM Location
WHERE opened_on >= '2000-01-01';

SELECT * FROM Location
WHERE opened_on BETWEEN '2000-01-01' AND '2003-01-01';


SELECT CURRENT_TIMESTAMP - opened_on AS "Days open"
FROM Location;
```



Subtracting a **timestamp** from a **timestamp** produces a value of type **interval**.

Defining Interval Fields

We can specify the resolution we desire.



```
ALTER TABLE Location
ADD COLUMN days_online interval day;

UPDATE Location SET days_online = 'P2Y1M1W1DT1H1M1S';
```

- **P** indicates the formatting used.
- **2Y1M1W1D** adds 2 years, 1 month, 1 week and 1 day.
- **T** indicates the following is referring to time
- And **1H1M1S** adds 1 hour, 1 minute and 1 second.

name	days_online
character vary	interval day
Location 4	2 years 1 mon 8 days
Location 5	2 years 1 mon 8 days
Location 6	2 years 1 mon 8 days
Location 2	2 years 1 mon 8 days
Location 3	2 years 1 mon 8 days
Headquarters	2 years 1 mon 8 days

We learned ...

- That enumerate types are data types that are ordered lists in nature and that we can use this to sort the records.
- That there is a data type called UUID to store and validate universal identifiers.
- That we can store JSON objects in a field and we can query them using specific SQL operators.
- That any type can be used as an array of any dimension.
- That we can store binary files directly into a field.
- That there are specific data types to manage time-related data.

A large group of people, mostly young adults, are posing for a group photo in a room with a projector screen in the background. They are arranged in several rows, with some people sitting on the floor in the front. Many are making peace signs or other celebratory gestures. The text "THANK YOU" is overlaid in large white letters in the center of the image.

THANK YOU

Contact Details
DCI Digital Career Institute gGmbH