

Digital Career Institute

Python Course - Advanced SQL



Goal of the Submodule

The goal of this submodule is to help the student understand how to work with advanced SQL queries and data types. By the end of this submodule, the learners will be able to understand:

- How to query multiple tables using different types of **JOIN** clauses.
- How do **JOIN** operations work on the background.
- How to use advanced SQL data types.
- How to group and use aggregate functions to extract information from the data tables.
- How to combine all this knowledge to define complex analytical SQL queries.

Topics

- How to query multiple tables.
- How to query multiple tables using different types of **JOIN** clauses.
- How to define and use advanced data types.
- How to group and calculate statistics on a set data tables.

Joining Multiple Tables

Introducing JOIN

DLI

```
SELECT fields  
FROM table_a  
JOIN table_b on conditions;
```

The **JOIN** clause takes all combinations of records between **table_a** and **table_b** and returns those that match the indicated **conditions**.

The **conditions** **MUST** include the equality between the foreign key and primary key of those tables.

Introducing JOIN

DLI

```
SELECT *  
FROM City  
JOIN Country ON Country.id = City.country_id;
```

Country	id	name
	1	Germany
	2	France
	3	USA

City	name	country_id
	Marseille	2
	Chicago	3
	Berlin	1

SQL Query

Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany

Implicit JOIN

DLI

Non implicit

Modern, default
joining syntax.

```
SELECT City.name, Country.name  
FROM City  
JOIN Country ON Country.id = City.country_id;
```

=

Implicit

Older, but still
used sometimes.

```
SELECT City.name, Country.name  
FROM City, Country  
WHERE Country.id = City.country_id;
```

UPDATE Across Tables

Implicit

```
UPDATE City SET country = Country.name  
FROM Country WHERE City.country_id = Country.id;
```

*In some RDBMS, like PostgreSQL, an **UPDATE** across tables is still done with an implicit join.*

JOIN & Order

DLI

```
SELECT *  
FROM City  
JOIN Country on Country.id = City.country_id  
ORDER BY Country.name;
```

Country	id	name
	1	Germany
	2	France
	3	USA

City	name	country_id
	Marseille	2
	Chicago	3
	Berlin	1

SQL Query

Result			
name	country_id	id	name
Marseille	2	2	France
Berlin	1	1	Germany
Chicago	3	3	USA

Types of JOIN

DLI

But what if ...

Country	id	name
	1	Germany
	2	France
	3	USA
	4	Spain



There are countries not being used in our city table.

City	name	country_id
	Marseille	2
	Chicago	3
	Berlin	1
	Barcelona	
	Salzburg	7



There are cities with **null** country_id or an id that does not exist.

Types of JOIN

The result is still the same.

The results include only those combinations of records that match the **on** condition.


Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany

*This type of join is called
INNER JOIN.*

Types of JOIN

INNER JOIN

DLI



The **JOIN** clause is a shortcut for the **INNER JOIN** clause.

For clarity purposes, it is usually preferred to use **INNER JOIN**.

```
SELECT {fields}  
FROM {table_a}  
INNER JOIN {table_b} on {conditions};
```

Other Types of JOIN

Very often, there is a main table and a secondary table.

Losing all the records in the main table whose foreign key is not found on the secondary table may not be desirable.

Very often, the project requires to create this kind of result.




Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany
Barcelona			
Salzburg	7		

*This type of join is called
LEFT OUTER JOIN or,
simply, **LEFT JOIN**.*

LEFT OUTER JOIN

DLI



LEFT OUTER JOIN

```
SELECT *  
FROM City  
LEFT JOIN Country on Country.id = City.country_id;
```

Records **matching** the on
condition


+

Records from the main table
without a match.

Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany
Barcelona			
Salzburg	7		

RIGHT OUTER JOIN

DLI



RIGHT OUTER JOIN

```
SELECT *  
FROM City  
RIGHT JOIN Country on Country.id = City.country_id;
```

Records **matching** the on
condition


+

Records from the **secondary
table without a match.**

Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany
		4	Spain

FULL OUTER JOIN

DLI



FULL OUTER JOIN

```
SELECT *  
FROM Country  
FULL JOIN City on Country.city_id = city.id;
```

Records **matching** the on
condition

+


Records from **any table**
without a match.

Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany
Barcelona			
Salzburg	7		
		4	Spain




Summary

DLI

INNER



OUTER



Using JOINS

Filtering JOINS

```
SELECT Location.name, City.name FROM Location  
INNER JOIN City on City.id = Location.city_id  
WHERE City.country_id = 3;
```

Location	name	city_id
	Headquarters	2
	Location 2	1
	Location 3	2
	Location 4	3
	Location 5	4
	Location 6	

Result	
name	name
Headquarters	Chicago
Location 3	Chicago

Filtering JOINS

OK

```
SELECT Location.name, City.name FROM Location  
INNER JOIN City on City.id = Location.city_id  
WHERE City.country_id = 3;
```

These two statements are equivalent.

?

```
SELECT Location.name, City.name FROM Location  
INNER JOIN City  
on City.id = Location.city_id AND City.country_id = 3;
```

Filtering JOINS

```
SELECT {fields} FROM {table_a}  
INNER JOIN {table_b} on {join_conditions}  
WHERE {filter_conditions};
```

For the sake of semantics and readability
the **on** keyword should define only the
joining conditions and the **WHERE** clause
should define only the **filtering
conditions**.

Multiple JOINS

```
SELECT Location.name, City.name, Country.name FROM Location  
LEFT JOIN City ON City.id = Location.city_id  
LEFT JOIN Country ON Country.id = City.country_id;
```

Location	name	city_id
	Headquarters	2
	Location 2	1
	Location 3	2
	Location 4	3
	Location 5	4
	Location 6	

Result		
name	name	name
Headquarters	Chicago	USA
Location 2	Marseille	France
Location 3	Chicago	USA
Location 4	Berlin	Germany
Location 5	Barcelona	
Location 6		

Multiple JOINS

DLI

*The order of the **JOIN** clauses matters.*

OK

```
SELECT Location.name, City.name, Country.name FROM Location
LEFT JOIN City on City.id = Location.city_id
LEFT JOIN Country on Country.id = City.country_id;
```



```
SELECT Location.name, City.name, Country.name FROM Location
LEFT JOIN Country on Country.id = City.country_id
LEFT JOIN City on City.id = Location.city_id;
```


At this point **City** is still undefined.

How JOINS work

How JOINS Work

The basic **JOIN** operations can be visualized as Venn diagrams from set theory.

But the **JOIN** operation does not actually work like this and not all **JOIN** types can be explained this way.



Venn diagram of an INNER JOIN.

Remember the first slide on JOINS:



The **JOIN** clause **takes all combinations of records** between **table_a** and **table_b** and returns those that match the indicated **conditions**.

CROSS JOIN

```
SELECT *  
FROM City  
CROSS JOIN Country;
```

The results include **every possible combination** between records on both tables.

All **JOIN** operations derive from a **CROSS JOIN**. They are filters of this main cross table.

Result

	name character va	country_id integer	id integer	name character va
1	Marseille	2	1	Germany
2	Marseille	2	2	France
3	Marseille	2	3	USA
4	Marseille	2	4	Spain
5	Chicago	3	1	Germany
6	Chicago	3	2	France
7	Chicago	3	3	USA
8	Chicago	3	4	Spain
9	Berlin	1	1	Germany
10	Berlin	1	2	France
11	Berlin	1	3	USA
12	Berlin	1	4	Spain
13	Barcelona		1	Germany
14	Barcelona		2	France
15	Barcelona		3	USA
16	Barcelona		4	Spain
17	Salzburg	7	1	Germany
18	Salzburg	7	2	France
19	Salzburg	7	3	USA
20	Salzburg	7	4	Spain

CROSS JOIN

```
SELECT *  
FROM City  
CROSS JOIN Country;
```

=

```
SELECT *  
FROM City  
LEFT JOIN Country on True;
```

=

```
SELECT *  
FROM City  
INNER JOIN Country on True;
```

Result

	name character va	country_id integer	id integer	name character va
1	Marseille	2	1	Germany
2	Marseille	2	2	France
3	Marseille	2	3	USA
4	Marseille	2	4	Spain
5	Chicago	3	1	Germany
6	Chicago	3	2	France
7	Chicago	3	3	USA
8	Chicago	3	4	Spain
9	Berlin	1	1	Germany
10	Berlin	1	2	France
11	Berlin	1	3	USA
12	Berlin	1	4	Spain
13	Barcelona		1	Germany
14	Barcelona		2	France
15	Barcelona		3	USA
16	Barcelona		4	Spain
17	Salzburg	7	1	Germany
18	Salzburg	7	2	France
19	Salzburg	7	3	USA
20	Salzburg	7	4	Spain

INNER JOIN Example

DLI

```
SELECT Location.name, City.name FROM Location  
INNER JOIN City 1  
on City.id = Location.city_id 2  
WHERE City.country_id = 3; 3
```

1

CROSS JOIN

name	name
Location 2	Marseille
Location 3	Marseille
Location 4	Marseille
Location 5	Marseille
...	...

+



2

ON CONDITION

name	name
Headquarters	Chicago
Location 2	Marseille
Location 3	Chicago
Location 4	Berlin
Location 5	Barcelona



3

WHERE COND.

name	name
Headquarters	Chicago
Location 3	Chicago

INNER JOIN Example

DLI

```
SELECT *  
FROM City  
JOIN Country on Country.id = City.country_id;
```

Result			
name	country_id	id	name
Marseille	2	2	France
Chicago	3	3	USA
Berlin	1	1	Germany

CROSS JOIN

	name character va	country_id integer	id integer	name character va
1	Marseille	2	1	Germany
2	Marseille	2	2	France
3	Marseille	2	3	USA
4	Marseille	2	4	Spain
5	Chicago	3	1	Germany
6	Chicago	3	2	France
7	Chicago	3	3	USA
8	Chicago	3	4	Spain
9	Berlin	1	1	Germany
10	Berlin	1	2	France
11	Berlin	1	3	USA
12	Berlin	1	4	Spain
13	Barcelona		1	Germany
14	Barcelona		2	France
15	Barcelona		3	USA
16	Barcelona		4	Spain
17	Salzburg	7	1	Germany
18	Salzburg	7	2	France
19	Salzburg	7	3	USA
20	Salzburg	7	4	Spain

Edge Case Examples

The **on** condition can be anything that evaluates to **True** or **False**.

```
SELECT *
FROM City
JOIN Country on False;
```

or

```
SELECT *
FROM City
JOIN Country on
Country.id = City.country_id OR city_id + id = 5;
```

CROSS JOIN

	name character va	country_id integer	id integer	name character va
1	Marseille	2	1	Germany
2	Marseille	2	2	France
3	Marseille	2	3	USA
4	Marseille	2	4	Spain
5	Chicago	3	1	Germany
6	Chicago	3	2	France
7	Chicago	3	3	USA
8	Chicago	3	4	Spain
9	Berlin	1	1	Germany
10	Berlin	1	2	France
11	Berlin	1	3	USA
12	Berlin	1	4	Spain
13	Barcelona		1	Germany
14	Barcelona		2	France
15	Barcelona		3	USA
16	Barcelona		4	Spain
17	Salzburg	7	1	Germany
18	Salzburg	7	2	France
19	Salzburg	7	3	USA
20	Salzburg	7	4	Spain

We learned ...

- How to cross reference data from multiple tables using JOIN clauses.
- That there are different types of JOIN clauses we can use.
- How to filter and use multiple JOINS in one statement.
- That the joining operation actually filters a complete cross reference of records to produce the result.

A large group of diverse people, including men and women of various ages, are posing for a group photo in a room. They are arranged in several rows, with some sitting on the floor in the front. Many are making peace signs or other celebratory gestures. In the background, there is a large projector screen and a whiteboard. The overall atmosphere is positive and energetic.

THANK YOU

Contact Details
DCI Digital Career Institute gGmbH